

Constrained Optimization: Exercise 1

Beatrice Macchia, Riccardo Bosio

Politecnico di Torino

Student ids: s245844, s291299

s245844@studenti.polito.it, s291299@studenti.polito.it

I. PROBLEM OVERVIEW

The problem consists of implementing the projected gradient method to solve $\min_{x \in R^n} \sum_{i=1}^n (\frac{1}{4}x_i^4 + \frac{1}{2}x_i^2 - x_i)$ s.t. $1 \leq x_i \leq 2 \ \forall i$. We solve this constrained optimization problem with $n = 10^4$ and $n = 10^6$, both using exact derivatives and using finite differences to approximate the gradient. Then, a comparison between the behaviors of the two implementations must be made, using these values for the increment h : $h = 10^{-k}\hat{x}$, $k = 2, 4, 6, 8, 10, 12$. \hat{x} is the point at which the derivatives have to be approximated.

II. PROPOSED APPROACH

To solve this constrained optimization problem we first create *function_to_optimize.m* to introduce the function that we have to optimize and in *gradient.m* we build the gradient function using exact derivatives. Then we notice that the function is separable [1], so we create *separated_function_to_optimize.m*, which is a function $F: R^n \rightarrow R^n$ with each component of the sum as a row.

$$F(x_1, x_2 \dots x_n) = \begin{bmatrix} g(x_1) \\ g(x_2) \\ \vdots \\ g(x_n) \end{bmatrix} \quad (1)$$

where $g(x) = \frac{1}{4}x^4 + \frac{1}{2}x^2 - x$. Calling f the function to optimize, we get $f(x) = 1_n^T \cdot F(x)$, with 1_n vector of n ones.

Then we notice that:

$$\begin{aligned} f(x + he_1) - f(x) &= 1_n^T \cdot F(x + he_1) - 1_n^T \cdot F(x) = \\ &= 1_n^T \cdot (F(x + he_1) - F(x)) = g(x_1 + h) - g(x_1). \end{aligned}$$

At this point we can create *finite_difference_gradient.m* to implement the finite differences to compute the gradient.

In particular we exploit the property just discovered to approximate the i -th component of the gradient in the forward method, considering only the increment of the i -th component of F :

$$\nabla f(x) \simeq \begin{bmatrix} \frac{g(x_1+h)-g(x_1)}{h} \\ \frac{g(x_2+h)-g(x_2)}{h} \\ \vdots \\ \frac{g(x_n+h)-g(x_n)}{h} \end{bmatrix} = \frac{F(x+h) - F(x)}{h} \quad (2)$$

With a similar approach we get that we can approximate the gradient in the centered method in the following way:

$$\nabla f(x) \simeq \frac{F(x+h) - F(x-h)}{2 \cdot h}$$

The feasible set is $X = \{x \in R^n | 1 \leq x_i \leq 2 \ \forall x_i\}$ and we need a projection function to implement the projected gradient method. Since the feasible set is in the form $X = \{x \in R^n | L \leq x_i \leq U \ \forall x_i\}$, we use the following projection:

$$[\pi_X(y)]_i = \begin{cases} y_i & L \leq y_i \leq U \\ L & y_i < L \\ U & y_i > U \end{cases}$$

This function is in *projection.m*. Then we can implement the projected gradient method in *projected_gradient_method.m*, with backtracking using the Armijo condition. Finally, in *C1.m* we can run the projected gradient method and make comparisons.

III. RESULTS

The projected gradient method that we have implemented always reaches the solution x_k equal to a vector of n ones at the k -th iteration. The corresponding value of the function is $f(x_k) = -2500$ for $n = 10^4$ and $f(x_k) = -250000$ for $n = 10^6$.

Backtracking iterations are always equal to zero for all the methods and for all values of h .

We tested our code with the following parameters:

Initializations	
Parameter	Value
alpha0	1
kmax	1000
tolgrad	1e-12
c1	1e-4
rho	0.8
btmax	50
gamma	0.1
tolx	1e-6

We chose $x_0 = (1.5, 1.5, \dots, 1.5)^T$, which is inside the feasible set, as starting point.

The results in terms of number of iterations and computing time are summarized in the following tables.

$n = 10^4$			
Number of iterations			
h	Exact derivatives	Finite differences (centered)	Finite differences (forward)
	3		
$10^{-2}\hat{x}$		2	2
$10^{-4}\hat{x}$		3	3
$10^{-6}\hat{x}$		3	3
$10^{-8}\hat{x}$		3	3
$10^{-10}\hat{x}$		3	3
$10^{-12}\hat{x}$		3	3
$n = 10^4$			
Computing time(seconds)			
h	Exact derivatives	Finite differences (centered)	Finite differences (forward)
	0.040461		
$10^{-2}\hat{x}$		0.034575	0.033844
$10^{-4}\hat{x}$		0.040235	0.037343
$10^{-6}\hat{x}$		0.037843	0.03349
$10^{-8}\hat{x}$		0.033699	0.033699
$10^{-10}\hat{x}$		0.035738	0.030167
$10^{-12}\hat{x}$		0.033545	0.028171
$n = 10^6$			
Number of iterations			
h	Exact derivatives	Finite differences (centered)	Finite differences (forward)
	3		
$10^{-2}\hat{x}$		2	2
$10^{-4}\hat{x}$		3	3
$10^{-6}\hat{x}$		3	3
$10^{-8}\hat{x}$		3	3
$10^{-10}\hat{x}$		3	3
$10^{-12}\hat{x}$		3	3
$n = 10^6$			
Computing time(seconds)			
h	Exact derivatives	Finite differences (centered)	Finite differences (forward)
	1.519397		
$10^{-2}\hat{x}$		2.827075	2.354431
$10^{-4}\hat{x}$		4.068107	3.396018
$10^{-6}\hat{x}$		3.985171	3.412225
$10^{-8}\hat{x}$		4.016851	3.539372
$10^{-10}\hat{x}$		3.834642	3.377349
$10^{-12}\hat{x}$		4.072665	3.463752

Also we tested our code changing the starting point from $x_0 = (1.5, 1.5, \dots, 1, 5)^T$ to $x_0 = (0, 0, \dots, 0)^T$, which is an external point with respect to the feasible set X . In terms

of iterations, with the new x_0 we got that all the methods need only 1 iteration both for $n = 10^4$ and $n = 10^6$. Talking about computing time, as for $x_0 = (1.5, 1.5, \dots, 1, 5)^T$, there are not significant differences with $n = 10^4$. Moreover for $n = 10^6$ the behavior is similar for both values of x_0 : the exact derivatives method is the more efficient and the forward one performs better among the finite differences methods.

IV. DISCUSSION

Looking at the results we can observe that the finite differences method with $h = 10^{-2}\hat{x}$ performs better in terms of number of iterations for both values of n . We get the result in 2 iterations, while in the other cases 3 iterations are needed.

Instead, in terms of computing time we can see that for $n = 10^4$ we get the fastest result using the forward finite differences method with $h = 10^{-12}\hat{x}$. However, comparisons between the computing times with $n = 10^4$ are not that meaningful, as they differ only in hundredths of a second.

For $n = 10^6$ we can note slightly more relevant differences in terms of computing time: in general the exact derivatives method performs better with respect to the other options.

V. APPENDIX

A. function_to_optimize.m

```

1 %Create the function which has to be optimized
2
3 function y = function_to_optimize(x,n)
4 %
5 %INPUTS:
6 %x=column vector of length n;
7 %n=number of dimensions of x;
8 %OUTPUTS:
9 %y=real scalar value equal to f(x)
10 %
11 y=0;
12 for i=1:n
13     y=y+(1/4*x(i)^4+1/2*x(i)^2-x(i));
14 end
15 end

```

B. separated_function_to_optimize.m

```

1 function Y = separated_function_to_optimize(x,n)
2 %
3 %INPUTS:
4 %x=column vector of length n;
5 %n=number of dimensions of x;
6 %OUTPUTS:
7 %Y=separated function R^n->R^n.
8 %
9 Y=zeros(n,1);
10 for i=1:n
11     Y(i)=1/4*x(i)^4+1/2*x(i)^2-x(i);
12 end
13
14 end

```

C. gradient.m

```
1 %Gradient function of the function to be ...
   optimized
2 function gradfx=gradient(x,n)
3 %
4 %INPUTS:
5 %x=column vector of size n;
6 %n=number of dimensions;
7 %
8 %OUTPUTS:
9 %gradf=column vector where i-th component is ...
   the partial derivative of f
10 %with respect to the i-th component of x.
11 %
12 gradfx=zeros(n,1);
13 for i=1:n
14     gradfx(i)=x(i)^3+x(i)-1;
15 end
16 end
```

D. finite_difference_gradient.m

```
1 %Compute the gradient with the finite ...
   difference method
2
3 function ...
   gradfx=finite_difference_gradient(F,x,n,k,type)
4 %
5 %function ...
   gradfx=finite_difference_gradient(f,x,n,h,type)
6 %
7 %INPUTS:
8 %F=separated function  $R^n \rightarrow R^n$ , of which we ...
   want to find the gradient;
9 %x=column vector of dimension n;
10 %n=dimension of x;
11 %k=integer between 2 and 12 used to find h, ...
   the parameter of the finite
12 %difference method;
13 %type='fw' or 'c' to choose the ...
   forward/centered finite difference method
14 %respectively.
15 %
16 %OUTPUTS:
17 %gradfx=column vector of dimension n which ...
   approximates the gradient of f
18 %in x.
19 %
20
21 gradfx=zeros(n,1);
22
23 h=10^(-k)*norm(x);
24
25 switch type
26     case 'fw'
27         xh=x+h;
28         gradfx=(F(xh)-F(x))/h;
29     case 'c'
30         xh_plus=x+h;
31         xh_minus=x-h;
32         gradfx=(F(xh_plus)-F(xh_minus))/(2*h);
33     otherwise %we do forward one
34         xh=x+h;
35         gradfx=(F(xh)-F(x))/h;
36 end
37 end
```

E. projection.m

```
1 %Function to project xk_bar into the ...
   feasible set
2 %
3 %
4 %
5 function x_hat=projection(x,lower,upper)
6 %
7 %function x_hat=projection(x,lower,upper)
8 %
9 %INPUTS:
10 %x is a column vector containing n components
11 %lower is a column vector where the i-th ...
   element is the lower bound of the
12 %i-th dimension of the feasible set
13 %upper is a column vector where the i-th ...
   element is the upper bound of the
14 %i-th dimension of the feasible set
15 %
16 %OUTPUT:
17 %x_hat is equal to x if x is already in the ...
   feasible set otherwise it
18 %is updated to the projection of x in ...
   the boundary of the feasible set
19 %
20 x_hat=max(min(x,upper),lower);
21 end
```

F. projected_gradient_method.m

```
1 %Projected gradient method function
2 function [xk, fk, gradfk_norm,  $\Delta xk\_norm$ , k, ...
   xseq, btseq]=...
   projected_gradient_method(x0, f, gradf, ...
   alpha0, kmax, tolgrad, c1,...
   rho, btmax, gamma, tolx, Pi_X)
3 %
4 %function [xk, fk, gradfk_norm,  $\Delta xk\_norm$ , k, ...
   xseq, btseq]=
   projected_gradient_method(x0, f, gradf, ...
   alpha0, kmax, tolgrad,
   c1, rho, btmax, gamma, tolx, Pi_X)
5 %
6 %INPUTS:
7 %x0=n-dimensional column vector;
8 %f=function  $R^n \rightarrow R$ ;
9 %gradf=function that describes the gradient ...
   of f;
10 %alpha0=first constant that multiplies the ...
   descent direction;
11 %kmax=maximum number of iterations;
12 %tolgrad=tolerance for the gradient (used as ...
   a stopping criterion);
13 %c1=Armijo condition factor that must be a ...
   scalar in (0,1);
14 %rho=fixed factor lesser than one used for ...
   reducing alpha0 at each
15 %iteration;
16 %btmax=maximum number of steps of the ...
   backtracking strategy to update
17 %alpha;
18 %gamma=constant that multiplies the descent ...
   direction before projection;
19 %tolx=tolerance for the norm of the ...
   difference between two consecutive
20 %xk (used as a stopping criterion);
21 %Pi_X=function that handles the projection;
22 %
23 %OUTPUTS:
```

```

28 %xk=the last x computed by the function;
29 %fk=f(xk);
30 %gradfk_norm=norm of gradf(xk);
31 %k=last iteration performed;
32 %xseq=n-by-k matrix where the columns are ...
    the xk computed during the
33 %iterations;
34 %btseq=1-by-k vector where elements are the ...
    number of backtracking
35 %iterations at each optimization step.
36 %
37
38 %Initializations
39 xseq=zeros(length(x0),kmax);
40 btseq=zeros(1,kmax);
41 xk=Pi_X(x0);
42 fk=f(xk);
43 k=0;
44 gradfk_norm=norm(gradf(xk));
45 Δxk_norm=tolx+1; %to be sure of entering the ...
    while
46
47 %create function to handle armijo conditions
48 f_armijo=@(fk,alpha,xk,pk) ...
    fk+c1*alpha*gradf(xk)*pk;
49
50 while k<kmax && gradfk_norm>tolgrad && Δ...
    xk_norm>tolx
51     %Find the descent direction
52     pk=-gradf(xk);
53     xk_bar=xk+gamma*pk;
54     xk_hat=Pi_X(xk_bar);
55
56     %Reset alpha
57     alpha=alpha0;
58
59     %Compute the new xk
60     pik=xk_hat-xk;
61     xnew=xk+alpha*pik;
62
63     %Compute f in xnew
64     fnew=f(xnew);
65
66     bt=0;
67
68     %Start backtracking
69     while bt<btmax && ...
        fnew>f_armijo(fk,alpha,xk,pik)
70         %Reduce alpha
71         alpha=rho*alpha;
72
73         %Update fnew and xnew
74         xnew=xk+alpha*pik;
75         fnew=f(xnew);
76
77         bt=bt+1;
78     end
79
80     %Update xk, fk, gradfk_norm, Δxk_norm
81     Δxk_norm=norm(xnew-xk);
82     xk=xnew;
83     fk=fnew;
84     gradfk_norm=norm(gradf(xk));
85
86     %Increase iteration
87     k=k+1;
88
89     %Store xk in k-th column of xseq
90     xseq(:,k)=xk;
91
92     %Store bt in btseq
93     btseq(k)=bt;
94 end
95

```

```

96 %Resize xseq and btseq
97 xseq=xseq(:,1:k);
98 btseq=btseq(1:k);
99
100 end

```

G. C1.m

```

1 clear all;
2 %
3 %
4 %Constrained optimization
5 %1)Projected gradient method
6 %
7 %
8 %Select a value for n
9 n=10^6; %We run C1 also with n=10^4
10
11 %Set parameters for the projected gradient ...
    method
12 x0=1.5*ones(n,1); %starting point
13 alpha0=1;
14 kmax=1000;
15 tolgrad=1e-12;
16 c1=1e-4;
17 rho=0.8;
18 btmax=50;
19 gamma=0.1;
20 tolx=1e-6;
21
22 %Create the function which has to be optimized
23 f=@(x) function_to_optimize(x,n);
24
25 %Create the separated function which has to ...
    be optimized
26 F=@(x) separated_function_to_optimize(x,n);
27
28 %Create the projection function
29 Pi_X=@(x) projection(x,ones(n,1),2*ones(n,1));
30
31 %Run the projected gradient method using the ...
    exact derivatives
32 %Create the gradient function
33 gradf=@(x) gradient(x,n);
34
35 tic
36 [xk, fk, gradfk_norm, Δxk_norm, k, xseq, ...
    btseq]=...
37     projected_gradient_method(x0, f, gradf, ...
        alpha0, kmax, tolgrad, c1,...
38     rho, btmax, gamma, tolx, Pi_X);
39 toc
40
41 disp('Number of iterations with exact ...
    derivatives:')
42 disp(k)
43
44 disp('Value of the function in xk:')
45 disp(fk)
46
47 %We have to try for different values of h
48 for j=2:2:12
49     j %print j
50
51     %Create the gradient computed with the ...
        finite difference method
52     tic
53     fd_gradf_c=@(x) ...
        finite_difference_gradient(F,x,n,j,'c'); ...
        %centered
54     [xk, fk, gradfk_norm, Δxk_norm, k, xseq, ...
        btseq]=...

```

```

55     projected_gradient_method(x0, f, ...
56         fd_gradf_c, alpha0, kmax, ...
57         tolgrad, c1, rho, btmax, gamma, ...
58         tolX, Pi_X);
59     toc
60     disp('Number of iterations with centered ...
61         finite differences:')
62     disp(k)
63     disp('Value of the function in xk:')
64     disp(fk)
65     tic
66     fd_gradf_fw=@(x) ...
67         finite_difference_gradient(F,x,n,j,'fw'); ...
68         %forward
69     [xk, fk, gradfk_norm, Δxk_norm, k, xseq, ...
70         btseq]=...
71         projected_gradient_method(x0, f, ...
72             fd_gradf_fw, alpha0, kmax, ...
73             tolgrad, c1, rho, btmax, gamma, ...
74             tolX, Pi_X);
75     toc
76     disp('Number of iterations with forward ...
77         finite differences:')
78     disp(k)
79     disp('Value of the function in xk:')
80     disp(fk)
81 end

```

REFERENCES

- [1] S. J.Nocedal, *Numerical Optimization*. Springer, 1999.